THE ENGINE OF NEURAL NETWORKS: GRADIENT-BASED OPTIMISATION

Training loop

1. Draw a batch of training samples x and corresponding targets y

2. **Run the network on x (this is called "forward pass"), obtain predictions y_pred**

3. **Compute the "loss" of the network on the batch, a measure of the mismatch between y_pred and y**

4. Update all weights of the network in a way that slightly reduces the loss on this batch.